

11 · Scorable negotiation & rational oracles

Siddharth M. Bhatia · siddharth@sidmb.com

July 23, 2026

Contents

The game: issues, score sheets, GameSpec	1
Swapping the game: presets	2
Exact solutions & descriptors	4
Rational oracles	5
Playing it: ScorableNegotiation	6
Rational agents: PolicyParticipant + the strategy zoo	7
Measuring divergence	8
Transcripts & Inspect	8

`interlens.arena.negotiation` is a **multi-party, multi-issue scorable negotiation** stack built for measuring how far an agent (LLM or rule-based) is from *game-theoretically optimal* play — not just whether it closed a deal. It ships a solver-verified game generator, exact normative benchmarks (Pareto frontier, Nash / Kalai–Smorodinsky / egalitarian / utilitarian / max-Nash-welfare points), a swappable **rational-oracle** layer that scores every move, the `ScorableNegotiation` scenario (structured binding offers separated from cheap talk), and a pool of pure-Python rational agents. Every generated instance carries its exact analysis, so per-turn regret and outcome-distance metrics are *exact*, not sampled.

This page is the negotiation-specific companion to [10 · Arena](#): the game lives in `arena/negotiation/`, the oracle types in `arena/oracles.py`, the scenario in `arena/scenarios/scorable.py`, and the rational agents in `PolicyParticipant`.

The game: issues, score sheets, GameSpec

A deal picks one option per issue and is a `Deal = tuple[int, ...]` of option indices. Each party holds a private additive `ScoreSheet` (a value per option,

plus a threshold τ — its BATNA); the analysis object is always the **surplus**
 $x_i(d) = u_i(d) - \tau_i$.

```
from interlens.arena.negotiation.space import Issue, DealSpace
from interlens.arena.negotiation.sheets import ScoreSheet, GameSpec

space = DealSpace((Issue("Site", ("North", "South")),
                    Issue("Fund", ("None", "1M", "5M"))))
alice = ScoreSheet("Alice", ((10, 0), (0, 3, 6)), threshold=5.0)
bob   = ScoreSheet("Bob",   ((0, 10), (6, 3, 0)), threshold=5.0)

alice.utility((1, 2))      # 0 + 6 = 6.0    (South, 5M)
alice.surplus((1, 2))      # 6 - 5 = 1.0

spec = GameSpec(space=space, sheets=(alice, bob),
                rounds=4, info="full", chat=True,      # protocol
                arms (DESIGN §3)
                proposer=0, veto=1, min_accept=None,   # agreement
                rule (None = unanimity)
                discount=1.0, breakdown_risk=0.0)      #
                impatience knobs the oracles read
spec.to_json()      # round-trips via GameSpec.from_json
                    - drops into Instance.payload
```

`GameSpec.surplus_matrix()` is the dense $|D| \times n$ surplus array every solution concept consumes; `feasible_mask()` is the boolean mask of deals that pass the agreement rule (proposer + veto clear, and $\geq \text{min_accept}$ parties clear their threshold).

Swapping the game: presets

The situation is a **swappable axis**, like the oracle stack — “two agents play the ultimatum game” is one call. `arena/negotiation/games.py` is a registry mapping a name to a literature-grounded **preset**: a factory returning (`GameSpec`, `analysis`, `protocol_cfg`), the game plus its exact analysis plus the scenario knobs that turn the one `ScorableNegotiation` engine into that game. Presets are *parameterizations, not new engines* — the same deal space, oracle stack, annotation, and atlas run unchanged on any of them.

```
from interlens.arena.negotiation import games
```

```

games.PRESETS #
    {'scorable', 'ultimatum',
     'divide_dollar', 'bilateral_multiissue'}
game, analysis, protocol_cfg = games.make_preset("ultimatum",
    pie=10, n_options=11)
game.n_parties, analysis["deal_space_size"] # 2, 11 (one
    split per option)
protocol_cfg #
    {'single_shot': True, 'fixed_proposer': True}
# the arena bridge: a solver-verified Instance + the cfg to play it
# under (reuses generate.build_instance)
instance,
    protocol_cfg = games.build_preset_instance("divide_dollar",
    n_parties=3, rule="majority")

```

Preset	Game	Rational anchor
scorable (default)	the multi-party, multi-issue repaired game (\$2 generator; a thin alias, defaults unchanged)	the \$4 oracle stack
ultimatum	take-it-or-leave-it N=2 pie split (J=1, $\tau=0$, rounds=1, single-shot responder-only vote)	SPE: proposer keeps the pie, responder accepts any positive share (Güth et al. 1982 human-rejection contrast; Rubinstein 1982 §5)
divide_dollar	N-party discrete-shares split, rotating proposer, multi-round, unanimity/majority	Baron–Ferejohn 1989 / Okada 1996 ($v_i = 1/n$ — the equilibrium oracle's own anchor)
bilateral_multiissue	DoND-style N=2, J-issue, private values	Lewis et al. 2017 lineage

Preset	Game	Rational anchor
	(the generator at n_parties=2)	

`protocol_cfg` is passed as the scenario `cfg` (or merged into it) so the shared state machine runs the preset's protocol variant. `single_shot` skips the round-robin and goes straight to the propose→vote forced final; `fixed_proposer` pins the opener seat (no rotation) — both default off, so the standard multi-round game is unchanged. Majority vs unanimity is just `GameSpec.min_accept`. The experiment runner exposes all four as `run.py --game {scorable, ultimatum, divide_dollar, bilateral_multiissue}` (+ `--game-arg KEY=VALUE` preset knobs).

Exact solutions & descriptors

`solutions.analyze(space, sheets)` returns the full, JSON-serializable analysis every generated instance ships with — the normative benchmarks and the score-sheet descriptors that tell you whether the game is actually contestable.

```
from interlens.arena.negotiation import solutions
an = solutions.analyze(space, (alice, bob))
an["deal_space_size"]          # 6
an["ir_count"], an["pareto_count"] # |IR set|, |Pareto frontier|
an["ir_pareto_fraction"]       # |IR ∩ Pareto| / |IR| — how
    near-zero-sum the acceptable set is
an["dominated_acceptable_fraction"] # 1 - that: the score-sheet-
    repair target (leave dominated deals in play)
an["sparsity"], an["pairwise_iou"] # TMLR descriptors: zero-
    option sparsity, pairwise score overlap
an["ideal_surplus"]            # per-party best feasible
    surplus b_i (the distance normalizer)
an["solutions"]["nash"]        # {"deal", "named",
    "surpluses", "utilities", "ties", "scale_invariant", ...}
```

The five solution concepts are `nash`, `kalai_smorodinsky`, `egalitarian`, `utilitarian`, `max_nash_welfare`; `all_solutions(space, sheets)` returns them as `SolutionPoints` (each

carries `.deal`, `.surpluses`, `.index`, `.ties`, `.scale_invariant`). Use **nash** / **kalai_smorodinsky** as the scale-invariant anchors — egalitarian and utilitarian are flagged `scale_invariant=False` (meaningful only on normalized surpluses).

Per-move divergence is a distance in scale-invariant normalized-surplus space:

```
U, tau = spec.utility_matrix(), spec.thresholds
idx = space.index_of((1, 2))
solutions.distance_to_frontier(U, tau, idx)           # 0.0
    iff the deal is Pareto-efficient
ks = solutions.all_solutions(space, (alice, bob))
    ["kalai_smorodinsky"]
solutions.distance_to_solution(U, tau, idx, ks.index) # 0.0
    iff the deal equals the KS point
```

Rational oracles

An Oracle scores every action a seat could take at a decision point and names the best one; the seat's **regret** is `value(best) - value(chosen)` in the game's value units — the centipawn-loss analog for negotiation.

```
from interlens.arena import Oracle, OracleVerdict, OracleRecord,
    Propose, Accept
verdict = OracleVerdict(
    action_values={Accept("01"): 6.0, Propose((1, 2)): 8.0}, #
    keyed by the typed action
    best=Propose((1, 2)),
    beliefs=None, # a
    belief oracle fills this with its posterior
    flags=[] #
    named hard-violation markers, e.g. "ir_violation"
    verdict.best_value() # 8.0
    verdict.divergence(Accept("01")) # 8.0 - 6.0 = 2.0 (regret of the
    chosen action, ≥ 0)
    verdict.extra # free-form per-verdict
    diagnostics: best-response surplus_loss, reservation, v*, ...
```

The concrete negotiation oracles live in `arena/negotiation/` (`beliefs.py` exact Bayes over an enumerated type grid, `acceptance.py` optimal-stopping accept/reject, `bestresponse.py` exact expectimax, `equilibrium.py` a Banks–Duggan stationary equilibrium) and subclass this generic ABC. An **inline annotation** scores a seat’s actual move against the oracle’s best on the same state and is persisted as an `OracleRecord`:

```
rec = OracleRecord.annotation(verdict, round=2, seat="Alice",
                              oracle="bestresponse",
                              chosen_action=Accept("01"),
                              turn_idx=5)
rec.divergence                # 2.0 (best_value - chosen_value) -
    the per-turn regret
rec.to_json()                 # lands in Episode.round_checkpoints
    (identified by a "verdict" key)
```

Playing it: ScorableNegotiation

`ScorableNegotiation` (`arena/scenarios/scorable.py`) plays the game through any participant with **harness-enforced channel separation**: each turn is a private scratchpad, an optional public message (cheap talk), and one binding formal move — Propose a complete deal (the registry stamps it P1, P2, ...), Accept / Reject a live offer by id, or Walk. Offers are addressed by id so “I accept” is never ambiguous with two proposals on the table.

```
import asyncio
from interlens import APIParticipant, UsageMeter
from interlens.arena import EpisodePool, EpisodeStore
from interlens.arena.scenarios import ScorableNegotiation

scenario = ScorableNegotiation()
instance = scenario.generate_instance(level=0, seed=1) # payload
    carries the GameSpec, solution = the analyze() dict
meter = UsageMeter(budget=25.0)
player = APIParticipant(name="player", model_id="claude-sonnet-5",
    meter=meter, turn_token_floor=2048)

pool = EpisodePool(EpisodeStore("episodes/"), meter=meter)
episode = asyncio.run(pool.run_episode(scenario, instance,
    arm="moves_chat", participant=player))
```

```
episode.outcome      # {"success", "deal": bool, "deal_named":
    {...}, "usw"/"esw"/"nsw"/"gini",
    # "per_party_surplus", "ir_violations",
    "ceiling_surplus", "offers", ...}
```

Arms: `moves_chat` (binding moves + cheap talk), `moves_only` (the $\pm$ communication treatment), `team`, and `solo` (the communication-free control that must not match the multi-agent rate). Situational knobs go through `cfg` (rounds, personas, history window, `show_own_scores`) exactly as in 10 · [Arena](#). Turns persist as `TurnRecords` with `parsed_action` (`atype` / `deal_named` / `offer`); per-turn oracle annotations persist as `OracleRecord.to_json()` dicts in `Episode.round_checkpoints`.

Rational agents: `PolicyParticipant` + the strategy zoo

A `PolicyParticipant` mounts a pure-Python `policy(state) → Action` as a first-class participant, so a rational agent plays the *same* protocol as an LLM (and can sit across the table from one). The strategy zoo in `arena/negotiation/strategies.py` provides the executable opponents: time-dependent Faratin curves (`TimeDependentPolicy.boulware(...)` concedes near the deadline, `.conceder(...)` concedes early), `MicroPolicy`, `NaiveTitForTatPolicy`, `ToughPolicy`, and a `BayesianRationalPolicy`, plus `AC-next/AC-combi` acceptance conditions.

```
from interlens.arena.negotiation.policy_participant import
    PolicyParticipant
from interlens.arena.negotiation.strategies
    import TimeDependentPolicy
alice_agent = PolicyParticipant(
    name="Alice", policy=TimeDependentPolicy.boulware(), seat=0,
    sheet=alice, space=space, deadline=spec.rounds,
    discount=spec.discount)
# hand alice_agent (and a policy or LLM for bob) to
    EpisodePool.run_episode as the seat participants
```

A rational-vs-rational rollout should land near the frontier in `info="full"` and show calibrated inefficiency in `info="private"` (Myerson–Satterthwaite) — the sanity check before putting LLMs in the seats.

Measuring divergence

The `experiments/rational_agents/` measurement layer reads stored `Episode + Instance` and computes the **divergence atlas**: outcome quality (Pareto / NBS / KS distance via the `solutions` functions above, welfare, Gini), per-turn oracle regret (from the `round_checkpoints OracleRecord` rows) with the Park et al. no-regret trend tests, concession-curve τ /CRI, and a 12-row failure taxonomy tiered from fully-mechanical (IR violation) to LLM-judge (tactic-style). `annotate.py` writes the per-turn annotations; `analysis/report.py` aggregates a run into per-model per-arm tables + plots. See that experiment's README for the pipeline.

Transcripts & Inspect

Every turn's exact rendered prompt is persisted: the engine records each `SeatRequest.view` into `TurnRecord.view` by default (`EpisodePool(record_views=False)` to disable), so a transcript is faithful even if prompt code later drifts. `arena/export.py` turns a stored run into human-readable transcripts — `export.export_run(epochs_dir, instances_dir, out_dir)` (or `python -m interlens.arena.export --epochs ... --instances ... --out ...`) writes a markdown + self-contained HTML page per episode (game setup, per-turn scratchpad/message/action/oracle-regret, outcome + welfare) plus a run index; it renders old episodes (no view) too.

And the scenario is registered as an Inspect task via the preset registry, so `inspect eval / inspect view` work out of the box:

```
inspect eval interlens.arena.inspect/scorable --model anthropic/
  claude-sonnet-5 -T game=ultimatum -T arm=moves_only
```